

Chang Hyun Park

chang.hyun.park@it.uu.se

<https://iamchanghyunpark.github.io/>

(cell) +46 73 244 1152

RESEARCH INTERESTS

Interested in architectural support for system softwares: both native and virtualized.

Interested in the virtual memory system and address translation, cache and memory systems. Looking for opportunities to expand research (virtual memory system & address translation) to emerging memory systems such as stacked DRAM and NVRAM and GPUs/Accelerators.

Recent interest in providing system security through the I/O path from the CPU to the I/O device.

EDUCATION

Uppsala University, Uppsala, Sweden - **Supervisor:** [David Black-Schaffer](#)

Post Doctoral Researcher

September 2019 ~ August 2021 (*Expected*)

KAIST, Daejeon, Korea - **Advisor:** [Professor Jaehyuk Huh](#)

Ph.D., in Computer Science

August 2019

M.S. in Computer Science

February 2015

Hanyang University, Seoul, Korea

B.S. in Computer Science and Engineering

Graduated Summa Cum Laude

February 2013

PUBLICATIONS

Published

Sanghoon Cha, Bokyeong Kim, **Chang Hyun Park**, Jaehyuk Huh, "**Morphable DRAM Cache Design for Hybrid Memory Systems**", accepted for ACM Transactions on Architecture and Code Optimization (ACM TACO) Vol 16, Issue 3, August 2019

Jungi Jeong, **Chang Hyun Park**, Jaehyuk Huh, Seungryul Maeng, "**Efficient Hardware-assisted Logging with Asynchronous and Direct Update for Persistent Memory**", The 51st Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), October 2018.

Jeongseob Ahn, **Chang Hyun Park**, Taekyung Heo, Jaehyuk Huh, "**Accelerating Critical OS Services in Virtualized Systems with Flexible Micro-sliced Cores**", The European Conference on Computer Systems (EuroSys), April 2018.

Chang Hyun Park, Taekyung Huh, Jungi Jeong, Jaehyuk Huh, "**Hybrid TLB Coalescing: Improving TLB Translation Coverage under Diverse Fragmented Memory Allocations**" The 44th International Symposium on Computer Architecture (ISCA), June 2017

Daehoon Kim, **Chang Hyun Park**, Hwanju Kim, and Jaehyuk Huh, "**Virtual Snooping Coherence for Multi-core Virtualized Systems**", IEEE Transactions on Parallel and Distributed Systems (TPDS), July 2016

Chang Hyun Park, Taekyung Heo, and Jaehyuk Huh, "**Efficient Synonym Filtering and Scalable Delayed Translation for Hybrid Virtual Caching**", the International Symposium on Computer Architecture (ISCA), June 2016

Jeongseob Ahn, **Chang Hyun Park**, and Jaehyuk Huh, "**Micro-sliced Virtual Processors to Hide the Effect of Discontinuous CPU Availability for Consolidated Systems**", The 47th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO), December 2014

Under Submission

Bokyeong Kim, Sanghoon Cha, **Chang Hyun Park**, Soojin Hwang, Jongse Park, Jaehyuk Huh, "Undisclosed title (Double Blind Review process)", Submitted to ICS 2020.

Chang Hyun Park, Sanghoon Cha, Bokyeong Kim, Youngjin Kwon, David Black-Schaffer, Jaehyuk Huh, "Undisclosed title (Double Blind Review process)", Submitted to ISCA 2020.

Hassan Muhammad, **Chang Hyun Park**, David Black-Schaffer, "Undisclosed title (Double blind Review process)", Submitted to SIGMETRICS 2020

RESEARCH EXPERIENCE (BY IMPORTANCE)

Providing a TLB with adjustable coverage for diverse system/application mapping requirements

- **Improve TLB performance for heterogenous memory systems (DRAM, Stacked DRAM, NVM, etc.) with varying characteristics**
- Proposed a Hardware-Software joint effort to provide a TLB with adjustable TLB coverage
- Prior work are not prepared for such heterogenous memory systems
 - Rise of non-uniform memory, memory heterogeneity (3D-stacked memory, Non-volatile memory, etc.) require differing memory mappings per system memory configuration.
 - For example NUMA does not work well with large pages (Baptiste ATC '14)
 - Inter/Intra application memory mappings show significant difference of memory mapping
- A TLB scheme that performs well across the diverse memory mapping is necessary.
- We propose a hardware-software hybrid TLB coalescing, where:
 - The HW offers a TLB that can provide adjustable TLB coverage using the contiguity information encoded into the page table by the System SW
 - The SW chooses the best TLB coverage for each application, configures the HW to abide by the coverage, and writes contiguity information onto the unused portions of the page table entry
- **Published and Presented at ISCA 2017**

Speeding up virtual address translation whilst lowering power consumption

- **Improve Address translation scalability by delaying address translation to after cache miss, for the majority translations.**
- Proposed the use of virtually and physically addressed caches to reduce the majority of TLB lookups for every memory access to the cache hierarchy. Entire cache hierarchy is used as a hybrid addressed cache.
- Proposed a system which filters less common synonym causing memory accesses using simple filters. Every memory request accesses the filter prior to accessing cache or TLB.
 - For any synonym causing access, the TLB is looked up and cache is accessed in physical address (to prevent synonym problems).
 - For the majority of memory accesses which are synonym free, virtual address is used to access the cache hierarchy. (*No translation on cache hit*)
 - All cached blocks always use either the virtual or physical address. Coherence maintained using the address scheme for the block.
- Any last level cache misses of virtually addressed memory requests translate to physical address prior to accessing the main memory. A TLB like structure or proposed scalable translation structure buffer translations.
- Proposed system reduces power consumption and improves system performance.
- **Published and Presented at ISCA 2016.**

Providing Fault tolerance and Crash consistency to Storage Class Memory: Extending RAID

- **Leverage the fault tolerance of RAID to minimize logging write overhead in distributed NVM systems.**
- Made a case for the requirement of a system that provides fault tolerance and crash consistency, when employing storage class memory(SCM) with Non-volatile memories (NVM).
- Proposed adding fault tolerance to SCMs
 - NVM provides persistence, data is expected to be safe. However memory errors and failures result in possible loss of data.
 - Made a case of reusing RAID for SCMs to provide tolerance of random NVM failures.
- Proposed extending RAID abstraction to include crash consistency.
 - All NVM employing systems need to provide crash consistency, preventing any data inconsistency from a failure that occurs in the middle of a write.
 - Most trivial solution is journaling which backs up all data to be written, then updates the actual data.

- Currently filesystems provide crash consistency for disks; when memory is a storage class memory, a new journaling abstraction is required for ease of programming.
- NVMs have write endurance problems, thus proposed reducing journal writes via theorem.
- Made a case to use RAID technique in a distributed storage network over high performance network (Infiniband).
 - RAID allows distributing data and parity to different disks.
 - Proposed using the feature to distribute data to different nodes over Infiniband.
- Participated as a second author, Main developer of the SW system.
- **Research done during internship at Microsoft Research Asia, Summer 2016**

Proposing a HW logging mechanism that supports atomic durability for programmability and performance for transactions in NVMs

- Proposed a HW logging mechanism based on both redo logging.
- HW logging of transactions for NVMs provides three benefits
 1. Ease of programming: Programmer just specifies which data to log. HW does the logging.
 2. Performance benefits: HW approach optimizes the logging and does not pollute the d-cache or i-cache with unnecessary data/instructions.
 3. Provides crash consistency: Non-volatile memory (NVMs) require special attention to keep data consistent, even in the event of a power failure. To make sure data is recoverable during a power failure, crash consistency is required.
- The contribution of this work is that we propose a direct-asynchronous approach to logging. Direct meaning that the data is persisted into the actual data in the NVM from the cache hierarchy, and asynchronous meaning that the actual persisting of the data is not on the critical path of the transaction commit.
- To achieve such properties, we introduce a small DRAM-cache that buffers writes to prevent early write-backs, and we track updated data to flush when the transaction is being committed.
- Participated as a second author, providing valuable intuitions and advices.
- **Published in MICRO 2018**

Restricting cache coherence snooping domain for virtualized systems

- Snooping protocol has scalability and power issues in large systems. Virtualized systems isolate memory among virtual machines providing an opportunity to isolate snooping requests among virtual machine domains.
- Snoop requests only sent to physical cores which hosted the virtual CPUs of the current snoop domain (virtual machine domain).
- Studied the snooping effects on spatial scheduling of virtual CPUs: restricting relocation of virtual CPUs to subset of physical cores.
- Participated as a second author, providing valuable intuitions and advices.
- **Published in TPDS 2016.**

Improving execution of workloads suffering virtualization induced anomalies

- Introduced the virtual time discontinuity problem, which encompasses anomalies induced by virtualization such as lock preemptions or delayed interrupt problems.
- Proposed a holistic solution to the virtual time discontinuity problem by shortening the hypervisor scheduling time.
- Proposed a HW cache preservation and restoration mechanism to account for performance suffered by cache sensitive workloads which benefit from longer time slices.
- Participated as a second author, providing valuable intuitions and advices.
- **Published in MICRO 2014.**

Micro-slicing cores to service OS critical executions

- Extending from the virtualization induced anomalies, this work takes a pure software solution to the virtual time discontinuity problem.
- A few cores of a host system is selected as micro-sliced cores, where the time quantum of these cores are 100us, or 0.1ms (compared to the regular 30ms or 1ms).
- Whenever events that are caused due to virtual time discontinuity are triggered, the hypervisor inspects the vCPUs that trigger the events. If they are OS critical logic, these vCPUs are migrated to the micro-sliced cores to quickly execute past the OS critical logic.
- Participated as second author, and came up with the main idea of identifying the vCPUs that are executing critical OS logic when an event is triggered.
- **Published in EuroSys 2018.**

INDUSTRY EXPERIENCE

Microsoft Research Asia, Beijing

Summer of 2016

Proposed fault tolerance and crash consistency abstraction to Storage Class Memories
(See Research Experience)

Samsung Electronics, Korea

Summer of 2012

Refactored and shortened bootup code of Uboot bootloader used by Tizen based phones.

AWARDS & SCHOLARSHIPS

KAIST:

National Scholarship: Full scholarship

March 2013 - February 2019

Hanyang University:

Academic Scholarship: Full scholarship (3 semesters), Partial scholarship (3 semesters)

Honorable mention at the Capstone Design Fair 2012

Korea Foundation for Advanced Studies:

[EE and CS Graduate Student Scholarship](#)

April 2016 - February 2019

Microsoft Research Asia:

[Microsoft Research Asia PhD Fellowship 2017](#)

October 2017

Naver:

Naver PhD Fellowship 2017

October 2017

TEACHING EXPERIENCE

KAIST Teaching Assistant:

CS230 System Programming

CS311 Computer Organization

CS510 Graduate Computer Architecture

CC510 Introduction to Computer Application (Graduate, non-CS students)

TECHNICAL SKILLS

Programming Languages: C/C++, Python, Java

System Software:

Xen hypervisor: modified credit scheduler code and profiled vCPU scheduling

Linux kernel: modified memory management code

Simulators & Tools:

Marssx86: implemented address translation after last level cache misses

gem5: current undisclosed work is on implementing TLBs in gem5.

Intel Pintool: used to create an in house simulator of TLB and cache

REFERENCES

Dr. Jaehyuk Huh
Professor,
School of Computing,
KAIST
<http://calab.kaist.ac.kr/~jhuh/>
jhhuh@kaist.ac.kr

Dr. David Black-Schaffer
Professor,
Department of Information Technology
Uppsala University
<https://katalog.uu.se/profile/?id=N9-1278>
david.black-schaffer@it.uu.se